

Content pipeline folder architecture and operating guide

Robert and Leigh should work from \\Iggy-Nas\\Shared\\ContentPipeline\\output\\ias-linkedln. Open **CURRENT.md** for the latest completed set, read the plan and research brief, and save decisions in **Reviews**. The surrounding folders separate original private meeting material, generated content evidence, human decisions, preserved completed cycles and evaluation work.

This guide describes the deployed system on **October 7, 2026**. It does not propose a folder migration. The browser copy is **Content-Pipeline-Architecture.html** and the printable copy is **Content-Pipeline-Architecture.pdf**, in the same starting folder. The canonical editable source is n8n/linkedin/Content-Pipeline-Architecture.md in rkigara-prog/iggii-ai-stack; shared HTML/PDF copies are generated from that source.

Folder design at a glance

\\Iggy-Nas\\Shared	
└─ Transcripts	ORIGINAL PRIVATE INPUTS
└─ Krisp-API	Canonical API TXT + JSON pairs
└─ Krisp	Legacy TXT + retained API shadows
└─ older meeting-named folders	Retained; outside the active selector
└─ ContentPipeline	
└─ output	Historical shared output namespace
└─ ias-linkedln	CURRENT PRODUCTION WORKING AREA
└─ START-HERE.md	Short starting instructions
└─ Content-Pipeline-Architecture.md / .html / .pdf	
└─ CURRENT.md	Links to last complete archived set
└─ current-status.json	Completion index and file hashes
└─ theme-list-*.txt	Sanitized intermediate theme artifact
└─ content-brief-*.md	Research intermediate and reading copy
└─ content-candidates-*.json	Evidence-bound planner input
└─ linkedln-editorial-plan-*.md / .json	Final cycle outputs
└─ privacy-status.json	Present privacy handoff
└─ candidate-status.json	Expected automated handoff; currently absent
└─ *.cjs	Service helpers, not editorial documents
└─ Reviews	HUMAN EDITS AND APPROVAL RECORDS
└─ Archive	
└─ ias-linkedln	
└─ UTC-date-content-hash/	Five copied outputs + manifest.json
└─ .pending-UUID/	Only if a copy failed; not a final archive
└─ older flat files	Retained historical archives
└─ older flat outputs / traces	LEGACY AND EVALUATION MATERIAL
└─ omc-krisp	Compatibility selector for old workflow
└─ ias-linkedln-acceptance	SYNTHETIC ACCEPTANCE AREA
└─ transcripts/Krisp-API and	transcripts/Krisp
└─ archive-check*	Synthetic output and archive fixtures
└─ ias-linkedln-real-privacy	PRIVATE REAL-MEETING EVALUATION AREA
└─ transcripts/Krisp-API and	transcripts/Krisp
└─ ias-orchestration	SYNTHETIC HANDOFF CHECKS
└─ fixtures	

The wildcard names in this tree identify file families, not extra folders. Production has **no separate intermediate-processing directory**: sanitized themes, researched candidates and completed plans share ias-linkedln, with distinct filenames and machine proofs establishing their stage. Most extracted themes, privacy decisions, search replies and ranking work exist inside an n8n execution rather than as durable folders. Do not use an evaluation folder to find production drafts.

Exact host and container mappings

The Windows name Iggy-Nas resolves to Unraid at 192.168.113.18. If name resolution fails, replace Iggy-Nas with 192.168.113.18 in any Windows path below. The SMB share is **Shared**; these are private LAN paths.

Windows share path	Unraid host path	n8n container path and access
\\Iggy-Nas\\Shared\\Transcripts	/mnt/user/Shared/Transcripts	/data/transcripts, read-only mount
\\Iggy-Nas\\Shared\\Transcripts\\Krisp-API	/mnt/user/Shared/Transcripts/Krisp-API	/data/transcripts/Krisp-API, read-only
\\Iggy-Nas\\Shared\\Transcripts\\Krisp	/mnt/user/Shared/Transcripts/Krisp	/data/transcripts/Krisp, read-only
\\Iggy-Nas\\Shared\\ContentPipeline\\output	/mnt/user/Shared/ContentPipeline/output	/data/output, read/write mount
\\Iggy-Nas\\Shared\\ContentPipeline\\output\\ias-linkedln	/mnt/user/Shared/ContentPipeline/output/ias-linkedln	/data/output/ias-linkedln
\\Iggy-Nas\\Shared\\ContentPipeline\\output\\Archive\\ias-linkedln	/mnt/user/Shared/ContentPipeline/output/Archive/ias-linkedln	/data/output/Archive/ias-linkedln
No path in Shared; service persistence	/mnt/app_pool/appdata/n8n/data	/home/node/.n8n; execution infrastructure and private operator records

The collector writes directly to the canonical host Krisp-API path through its own read/write bind mount. Its private service state is under /mnt/app_pool/appdata/omc-krisp-transcript-ingest, outside Shared/Transcripts. Neither credentials nor collector state belong in editorial folders. Old /volume1 paths refer to the retired Synology arrangement, not a current mount. There is no deployed Krisp-ai folder in this pipeline.

Folder responsibilities

In this table, **San** means the published sanitizer IASLinkedlnSan01, **Enr** means IASLinkedlnEnr01, and **Plan** means IASLinkedlnPlan1. “None active” means the current content chain does not read the folder; it does not declare the folder safe to delete. Robert and Leigh should follow the editing policy even where their accounts technically have write access.

Exact Windows share path	Purpose	Why it is separate	File creator	Consumer	Robert or Leigh editing policy
\\Iggy-Nas\Shared	Private SMB share entry	Existing share authorization precedes folder permissions	Unraid share configuration/operators	Samba serves authorized accounts	Use the named child locations; do not change share settings
\\Iggy-Nas\Shared\Transcripts	Private source namespace	Keeps original meetings outside generated content	Intake services and historical imports	San scans only the two named source roots below; OMC has separate intake	Do not edit, rename or move source material
\\Iggy-Nas\Shared\Transcripts\Krisp-API	Authoritative API TXT/JSON pairs	Preserves API identity, metadata and collector ownership of intake	omc-krisp-collector; repairs restored canonical copies	San selector; separately, OMC transcript intake	Neither reviewer should edit; a JSON sidecar is not a draft
\\Iggy-Nas\Shared\Transcripts\Krisp	Legacy transcripts and retained API shadows	Supports legacy intake while canonical API copies take priority	Google Drive Transcript Pull/rcclone and historical imports/moves	San fallback and duplicate checks; separate OMC legacy intake	Do not edit; unknown dates and duplicates are retained
\\Iggy-Nas\Shared\ContentPipeline	Container for content assets	Existing content namespace outside raw inputs	Historical setup/operators	Through its output child	No routine edits at this level
\\Iggy-Nas\Shared\ContentPipeline\output	Mounted output parent; older flat files	Existing deployment mount and pre-cutover namespace	Original workflows, evaluation runs and historical operators	Current stages use only explicit child paths	Browse for history; do not move or overwrite ambiguous flat files
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin	Current production artifacts, proofs and guides	Isolates the active chain from legacy wildcard outputs and test data	San, Enr, Plan and approved documentation deployment	Enr reads approved themes; Plan reads proven candidates; archive helper reads all five	Read generated artifacts; edit canonical guides through Git, not live helpers
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin\Reviews	Human requested edits and approval records	Prevents editorial changes from invalidating generated evidence hashes	Robert and Leigh	No active workflow; people read these records	Both may create/edit documents here

Exact Windows share path	Purpose	Why it is separate	File creator	Consumer	Robert or Leigh editing policy
\\Iggys-Nas\Shared\ContentPipeline\output\Archive	Historical archive parent	Existing archive location retained during repair	Original archival process, whose trigger is unrecovered	No active workflow reads the older flat archives	Treat old files as history; no cleanup without classification
\\Iggys-Nas\Shared\ContentPipeline\output\Archive\ias-linkedln	Verified copies of completed production sets	Gives review a stable copy while working files can change	Plan completion helper; initial controlled bootstrap	CURRENT.md links and human review; helper checks repeat copies	Read only for editorial work; Robert administers recovery
\\Iggys-Nas\Shared\ContentPipeline\output\omc-krisp	Versioned compatibility selector	Keeps the older inactive workflow's helper reference intact	Earlier migration/intake repair	Original inactive sanitizer Yjc03gS873IHEJPI; not active San	Do not edit; active selector is inside ias-linkedln
\\Iggys-Nas\Shared\ContentPipeline\output\ias-linkedln-acceptance	Synthetic migration/acceptance outputs	Keeps invented test material out of production selectors	Inactive/manual acceptance tools and operators	Test workflows/checks only	Operator-manage d; Leigh has no access; not editorial content
\\Iggys-Nas\Shared\ContentPipeline\output\ias-linkedln-acceptance\transcripts	Synthetic source staging	Allows the selector to be checked without real input	Acceptance fixtures	Acceptance selector	Neither reviewer should use it as a source
\\Iggys-Nas\Shared\ContentPipeline\output\ias-linkedln-acceptance\transcripts\Krisp-API	Synthetic API-shaped pairs	Exercises API handling	Acceptance fixture preparation	Acceptance selector	Operator only
\\Iggys-Nas\Shared\ContentPipeline\output\ias-linkedln-acceptance\transcripts\Krisp	Synthetic legacy-shaped TXT	Exercises fallback and duplicate rules	Acceptance fixture preparation	Acceptance selector	Operator only
\\Iggys-Nas\Shared\ContentPipeline\output\ias-linkedln-acceptance\archive-check	Initial synthetic writer/archive check	Retains check evidence without adding it to production	Archive acceptance tools	Temporary test workflow, now removed; helper checks	Operator only

Exact Windows share path	Purpose	Why it is separate	File creator	Consumer	Robert or Leigh editing policy
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check-final	Later synthetic archival check	Historical validation iteration; no permanent distinct role documented	Archive acceptance tools	Checks only	Operator only; do not mistake it for a final production cycle
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check-release	Final synthetic helper validation	Historical validation iteration	Archive acceptance tools	Checks only	Operator only
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-real-privacy	Real-meeting privacy evaluation results	Separates sensitive evaluation packets from production/research	Inactive IASPrivacyReal01 and private evaluation tools	Manual privacy evaluation, not active chain	Operator only; Leigh has no access
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-real-privacy\transcripts	Private evaluation source staging	Source context stays within its evaluation profile	Approved evaluation preparation	Evaluation selector	Neither reviewer should treat it as canonical intake
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-real-privacy\transcripts\Krisp-API	Staged API-shaped evaluation input	Mirrors API input structure	Approved evaluation preparation	Evaluation selector	Operator only
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-real-privacy\transcripts\Krisp	Staged legacy evaluation input	Mirrors legacy input structure	Approved evaluation preparation	Evaluation selector	Operator only
\\Iggy-Nas\Shared\ContentPipeline\output\ias-orchestration	Synthetic subworkflow handoff evidence	Tests waiting calls/failures without inference or production output	Orchestration fixture tools	Manual orchestration checks	Operator only; Leigh has no access
\\Iggy-Nas\Shared\ContentPipeline\output\ias-orchestration\fixtures	Synthetic themes/candidates/plans	Keeps fixture artifacts beneath their test profile	Fixture scripts/test workflows	Fixture stages only	Operator only

There are also seven retained meeting-named folders immediately under Transcripts. The active selector does not scan these; their earlier processing role is not documented sufficiently to justify consolidation. Their exact locations are:

Exact Windows share path	Purpose and separation	Creator	Active consumer	Editing policy
\\Iggy-Nas\Shared\Transcripts\chrome meeting - July 24, 2026 2-00-51 PM	Historical meeting import; original layout retained	Earlier import; exact writer unrecovered	None active	Do not edit/move

Exact Windows share path	Purpose and separation	Creator	Active consumer	Editing policy
\\Iggy-Nas\Shared\Transcripts\chrome meeting - July 24, 2026 9-30-03 AM	Historical meeting import; original layout retained	Earlier import; exact writer unrecovered	None active	Do not edit/move
\\Iggy-Nas\Shared\Transcripts\ms-teams meeting - July 24, 2026 3-10-38 PM	Historical meeting import; original layout retained	Earlier import; exact writer unrecovered	None active	Do not edit/move
\\Iggy-Nas\Shared\Transcripts\ms-teams meeting - July 24, 2026 11-01-51 AM	Historical meeting import; original layout retained	Earlier import; exact writer unrecovered	None active	Do not edit/move
\\Iggy-Nas\Shared\Transcripts\ms-teams meeting - July 24, 2026 1-28-51 PM	Historical meeting import; original layout retained	Earlier import; exact writer unrecovered	None active	Do not edit/move
\\Iggy-Nas\Shared\Transcripts\ms-teams meeting - July 24, 2026 9-03-08 AM	Historical meeting import; original layout retained	Earlier import; exact writer unrecovered	None active	Do not edit/move
\\Iggy-Nas\Shared\Transcripts\Zoom meeting - July 24, 2026 9-59-57 AM	Historical meeting import; original layout retained	Earlier import; exact writer unrecovered	None active	Do not edit/move

The fixture subfolders below are internal test counterparts, not another production architecture. Older flat synthetic copies also remain inside the initial archive-check/Archive. Every per-cycle hash directory inside these folders inherits the enclosing fixture's purpose and policy.

Exact Windows share path	Purpose	Why separate	File creator	Consumer	Editing policy
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check\output	Synthetic working files	Mirrors the tested helper profile; private fixture evidence	Archive acceptance tools	Fixture/check tools only	Operator only; neither reviewer uses it for content
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check\Archive	Synthetic archive parent	Mirrors the tested helper profile; private fixture evidence	Archive acceptance tools	Fixture/check tools only	Operator only; neither reviewer uses it for content
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check\Archive\output	Corrected synthetic archive namespace	Mirrors the tested helper profile; private fixture evidence	Archive acceptance tools	Fixture/check tools only	Operator only; neither reviewer uses it for content
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check-final\output	Synthetic working files	Mirrors the tested helper profile; private fixture evidence	Archive acceptance tools	Fixture/check tools only	Operator only; neither reviewer uses it for content

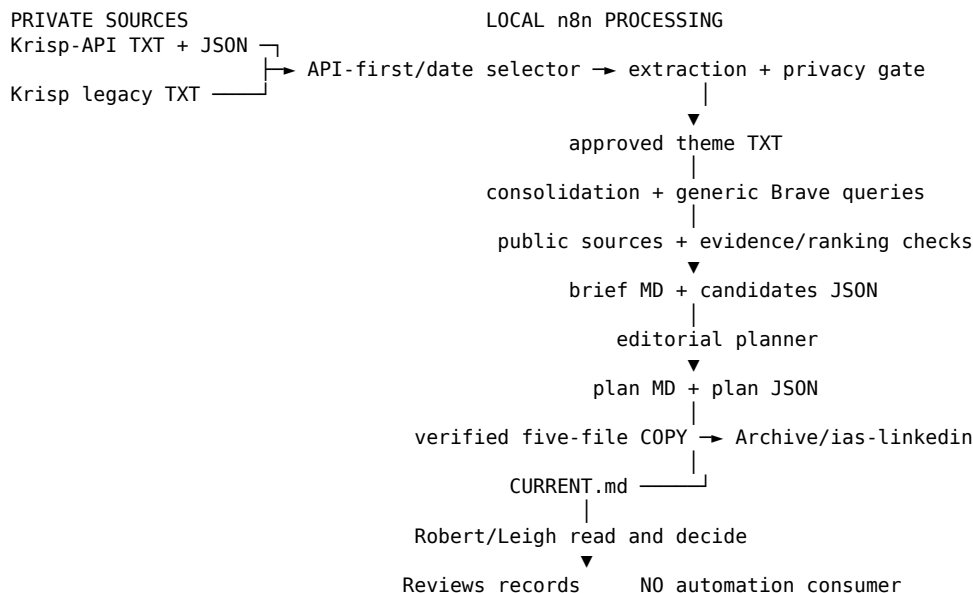
Exact Windows share path	Purpose	Why separate	File creator	Consumer	Editing policy
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check-final\Archive	Synthetic archive parent	Mirrors the tested helper profile; private fixture evidence	Archive acceptance tools	Fixture/check tools only	Operator only; neither reviewer uses it for content
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check-final\Archive\output	Corrected synthetic archive namespace	Mirrors the tested helper profile; private fixture evidence	Archive acceptance tools	Fixture/check tools only	Operator only; neither reviewer uses it for content
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check-release\output	Synthetic working files	Mirrors the tested helper profile; private fixture evidence	Archive acceptance tools	Fixture/check tools only	Operator only; neither reviewer uses it for content
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check-release\Archive	Synthetic archive parent	Mirrors the tested helper profile; private fixture evidence	Archive acceptance tools	Fixture/check tools only	Operator only; neither reviewer uses it for content
\\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin-acceptance\archive-check-release\Archive\output	Corrected synthetic archive namespace	Mirrors the tested helper profile; private fixture evidence	Archive acceptance tools	Fixture/check tools only	Operator only; neither reviewer uses it for content

For production cycle directories, the exact current Windows path is listed under Archiving and recovery. Each such child holds one content set, exists to preserve a distinct completed result, is created by Plan's completion helper, is read through CURRENT.md/human review, and should not be edited by either reviewer. A .pending-UUID child has the same creator and restricted access but is incomplete recovery material; no pending production directory was observed.

Service-only folders have no mapping within the Shared share:

Host folder	Purpose and separation	Creator	Consumer	Editing policy
/mnt/app_pool/appdata/n8n/data	n8n persistence/private operation records; separate from editorial data	n8n and authorized operator tools	n8n runtime and operator recovery	Neither reviewer edits it for content work
/mnt/app_pool/appdata/omc-krisp-transcript-ingest	Collector configuration/state/checkpoints; independent of original TXT/JSON source pairs	Collector and authorized operators	API intake service	Neither reviewer edits it for content work

Data flow and file authority



1. **Intake preserves originals.** The API collector creates a TXT transcript with its corresponding JSON metadata. The canonical pair is authoritative for API identity and started_at. Legacy TXT remains authoritative for meetings without an API pair. Restored API shadows in Krisp are retained copies, not an alternative canonical home. OMC's intake uses its own state/references; content processing does not rename meetings or change OMC identities.

2. **Selection establishes the input set.** San's select-transcripts.cjs scans Krisp-API first and Krisp second, recursively within those roots. It uses a rolling ten-day meeting-date window, not file modification time. API metadata must match the filename identity and contain a timezone-qualified date. Identical API shadows and matching legacy IDs are suppressed; conflicting API copies block selection. Undated legacy files stay on disk but are not selected. Copying a transcript to make it appear recent is not an intake method.

3. **Sanitization establishes research permission.** San extracts themes locally and runs a separate privacy review, then writes up to twenty anonymous theme bullets. privacy-status.json changes from pending to approved and binds the TXT's filename/hash to a run. Approval must still match the file and be no more than 24 hours old when downstream processing checks it. An old theme file alone is not authority to research; a new attempt invalidates the previous approval.

4. **Research creates generated evidence.** Enr reads the exact approved theme, consolidates topics, and issues generic queries through four Brave branches: theme-matched news, emerging/discovery news, recovery/industry news, and authoritative foundations. Raw transcripts are not web-search input. Source-family, evidence-fragment, scoring, corroboration and unsupported-claim rules separate qualified opportunities from the watchlist. The brief is the reading copy; the candidates JSON is the machine handoff. candidate-status.json binds candidate/brief hashes to the privacy proof.

5. **Planning creates an editorial proposal.** Plan verifies that handoff, selects from qualified candidates, preserves exact verified source URLs and writes both plan formats. Watchlist entries become research gaps, not approved post briefs. Webinar overlap still requires human review. A completed plan may have partial or insufficient_verified_topics status; technical completion does not certify editorial quality.

6. **Archiving establishes a completed record.** After both plan writes, the helper verifies the five-file set and publishes preserved copies plus a manifest. Only then does it replace current-status.json and CURRENT.md. These identify the latest successfully indexed set, not the newest file modification time. A failed cycle can leave newer partial working files while CURRENT.md correctly stays on the previous complete set.

7. **Review records human decisions.** Robert and Leigh create their own review documents. No active workflow reads Reviews or interprets “approved” as a trigger. There is no automated drafting, image generation or LinkedIn publishing action in this chain; automaticPublishingAllowed remains false.

Original inputs are authoritative for meeting evidence. JSON handoff proofs are authoritative for whether the next automated stage may consume generated files. The current bootstrapped set has privacy-status.json but no candidate-status.json: it is a preserved accepted set, not a currently resumable candidate handoff. Do not create that marker manually; Enr records a fresh proof during a successful automated cycle.

The candidates/plan JSON files are authoritative machine records for their stages; their Markdown partners are readable renderings. CURRENT.md and the archive manifest are authoritative for locating and verifying the completed set. Review records are authoritative only for the reviewers' stated decisions. None of these generated or historical copies replaces an original transcript.

The five output files

The current completed set has **October 6, 2026 filenames**. Its plan finished at **2026-10-07 01:53:26.887 UTC**, which was October 6 in New York. Filename dates are generated in America/New_York; archive directory dates use the plan's UTC generatedAt. Future cycles keep these basename patterns with their own dates.

Actual current filename	Contents	Intended use and automation role
theme-list-model-eval-qwen35-2026-10-06.txt	Up to twenty anonymous public-topic bullets derived from the privacy gate	San's generated intermediate; Enr reads it only with its matching privacy proof. Read for context; never restore private names or details
content-brief-model-eval-2026-10-06.md	Research opportunities, supporting sources/evidence, scoring context and watchlist items needing verification	Human research reading copy; archived and hash-bound alongside candidates. Read before approving an angle
content-candidates-model-eval-2026-10-06.json	Structured qualified theme-aligned/emerging opportunities, source/evidence records, diagnostics and watchlist	Enr's machine output and Plan's exact input. Do not hand-edit to promote a watchlist item or change evidence

Actual current filename	Contents	Intended use and automation role
linkedin-editorial-plan-model-eval-2026-10-06.md	Selected topic briefs, editorial angles, takeaways, verified URLs, webinar review and research priorities	Primary document for Leigh and Robert to read and discuss; save changes separately in Reviews
linkedin-editorial-plan-model-eval-2026-10-06.json	Structured planner decisions, selected/deferred topics, policy summary and validation warnings	Machine companion and archival lineage record; not a publishing queue or an approval form

The words model-eval, qwen35, and gpt-oss in inherited basenames/node labels do not make this current set a test run or establish the served model. The live stages use the migrated local home-chat route. Membership in ias-linkedln, published workflow bindings and matching completion proofs identify production. There are five content outputs; CURRENT.md, guides, handoff markers and the archive manifest are additional metadata, not extra content outputs.

Permissions and privacy boundaries

Identity	Production working area and archive	Reviews	Original transcripts and tests
Robert, SMB account rigarashi, UID 1000	Owner access; may administer files, but should not edit generated evidence during editorial review	Create/edit/read	Original source tree is accessible; current private test roots are owner-accessible
Leigh, SMB account leigh, UID 1001	Read/traverse; generated files and helpers are not writable	Create/edit/read	Source roots use the existing users group and are accessible; current test roots are mode 0700 and not accessible
n8n, UID/GID 1000, supplementary groups 100 and 1800	Creates/updates generated files, proofs, index and archives	Technically accessible, but no workflow consumes it	Read-only transcript mount; permitted private test access
API collector	Writes canonical API input through its specific bind mount	No content-review role	Docker has no configured non-root User override; input files are normalized to nobody:users. Private collector state stays outside the content output mount
OMC worker, UID 1002, users group 100	Excluded from the private production/archive child roots	Excluded	Existing source access supports separate OMC intake; identities and references remain unchanged
Guest/unauthorized share account	Shared SMB authorization denies access	Denied	No public access is configured

The private contentpipeline group, GID 1800, contains exactly Robert and Leigh. n8n has that supplementary GID separately. Production/archive directories are normally 2750 and generated files 0640; Reviews is 2770. Helpers explicitly assign GID 1800 because Unraid shfs did not reliably honor named ACLs or directory group inheritance. The group is restored by the boot hook; the Portainer n8n manifest persists its supplementary groups.

Folder separation is not a claim that raw transcripts are hidden from Leigh. Both reviewers currently have access to the private transcript source roots through Shared/users permissions. The important processing boundary is that San alone reads raw meetings and Enr's web queries use approved anonymous themes. Content artifacts are intended to support eventual public content, but they remain private working material, not automatically publishable data. The real-privacy evaluation area may itself contain sensitive staging/results and stays closed to Leigh. Existing same-model privacy-assurance limitations remain; a folder name or successful gate is not independent-model assurance.

Archiving and recovery

Confirmed current behavior: the active Plan completion node runs once after its two file writes, before returning success to Enr and San. It copies all five content outputs, verifies hashes and source lineage, then publishes a manifest and completion index. Archiving does **not** wait for human approval, move working outputs, archive review notes, remove transcripts, or expire old copies.

A finalized directory is named YYYY-MM-DD-`<64-character SHA-256 content hash>`. The date is the UTC plan completion date; the hash covers the ordered filenames, sizes and file hashes of the five-file set. The current exact directory is:

```
\\Iggly-Nas\Shared\ContentPipeline\output\Archive\ias-linkedln\2026-10-07-b8ec0ec5b54bb6bdf1f910c2c85f08afc7a23e54c289a1b7c3fef65dbd37b0a9
```

Exact Windows cycle folder	Purpose	Why separate	File creator	Consumer	Editing policy
\\Iggly-Nas\Shared\ContentPipeline\output\Archive\ias-linkedln\2026-10-07-b8ec0ec5b54bb6bdf1f910c2c85f08afc7a23e54c289a1b7c3fef65dbd37b0a9	Current preserved five-file set and checksum manifest	Keeps this distinct completed result stable across reruns	Controlled bootstrap of the already-accepted set; future directories are written by Plan's helper	CURRENT.md and human review	Neither reviewer edits the generated copy; Robert handles verified recovery

It contains the five files named above and manifest.json. Its manifest records bootstrap: true: the already-accepted set was copied during the repair, rather than by a new model run. That does not renew its privacy approval. The stored approval was made at 2026-10-07 01:49:45.270 UTC and expires at 2026-10-08 01:49:45.270 UTC if not invalidated sooner. The normal completion path still requires fresh, matching privacy/candidate proofs.

Situation	Current behavior and safe response
Another successful cycle on a different day	Creates another dated/hash directory and updates CURRENT.md

Situation	Current behavior and safe response
Changed output set on the same day	Creates a different hash directory; working basenames may be overwritten but prior finalized copies remain
Exact same five-file set archived again	Verifies the existing copy/manifest and reuses it; does not duplicate the set
Source changes, incomplete set, invalid path, or mismatching existing archive	Fails closed; does not silently overwrite the conflicting archive or delete sources
Interrupted copy	May leave .pending-UUID; it is not a completed cycle. Robert should inspect it with the private operation record and preserve it until its status is resolved
Archive/index failure after planning	Working files may exist, but parent success is blocked. Use the last valid CURRENT.md set; do not assume newer files form a completed cycle
Index repair	Robert/operator verifies finalized manifest hashes and lineage before restoring links. A controlled snapshot of an already-accepted set may rebuild an index; it is not permission to rerun research with an expired approval
Legacy flat Archive files	Remain historical. They are not automatically imported into the new manifest/index system

Copies are retained **indefinitely**, with no expiry or deletion job. The previous legacy archive trigger/age rule was not recoverable from deployed workflows, cron/User Scripts or retained migration records. The repair adopted completion-based copying and non-destructive retention; this should not be described as restoration of a proven old age threshold. Future expiry, storage limits, review-note retention and any automatic approval consumer remain unresolved design decisions requiring separate work. Completed archive directories are immutable by operating policy and read-only for Leigh, not a filesystem write-once guarantee against Robert or the service.

The authenticated SMB/copy-preservation gate and inactive synthetic n8n archive handoff passed in the earlier repair. The first Monday production run with the hook and a host reboot have not yet occurred. This guide's checks read paths, definitions and links only; they do not repeat pipeline execution or archive operations.

Leigh's practical walkthrough

1. Open \\Iggy-Nas\\Shared\\ContentPipeline\\output\\ias-linkedln. Sign in as leigh. If an old window still reports access denied, disconnect and reconnect Shared to refresh the repaired group membership. Use the IP-based path if the hostname does not resolve.
2. Open **Content-Pipeline-Architecture.html** in a browser, or its PDF for printing. START-HERE.md links to both. For work on a cycle, open **CURRENT.md** and check its completion time. This points to preserved copies, so the next scheduled run cannot replace the files you are reviewing.

3. Open the linked **linkedin-editorial-plan-model-eval-DATE.md** first, then **content-brief-model-eval-DATE.md** to inspect the evidence and cited sources. Markdown is plain text; a Markdown viewer improves formatting but is not required. JSON files are for detailed structured evidence, not documents to rewrite. Read the theme TXT only for topic context.
4. Create a file under \\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedin\Reviews, for example 2026-10-07-review.md or a Word document. Record the cycle directory/hash, exact source filenames, topic, requested wording changes, source-check result, webinar-overlap decision, reviewer/date and approval or rejection. A review-file date alone is not enough to distinguish same-day reruns.
5. Save proposed text and approvals in that review file. Do not replace the generated plan/JSON or copy private transcript details into a content draft. An “approved” note is a human record only: **automation does not consume review files**, start drafting, or post to LinkedIn.
6. For an older cycle, open \\Iggy-Nas\Shared\ContentPipeline\output\Archive\ias-linkedin. Choose a finalized date/hash directory, check its manifest/completion time and open its plan and brief. Ignore .pending-*. If several sets share a date, use their different hashes and your review's recorded cycle reference. Very old flat archives are one level higher at Archive and may not have manifests.
7. If CURRENT.md is older than expected, report the timestamp to Robert. Do not promote a file from a test folder or trigger an ad hoc rerun to make it look current. Robert can inspect the failed stage while review continues on the last complete set.

Design assessment

Separation or leftover	Assessment	Future consolidation consideration
Original Transcripts versus ContentPipeline	Justified: sources retain identities/metadata and are mounted read-only into n8n; downstream web research consumes anonymized derivatives	Keep the source/content boundary
Krisp-API versus Krisp	Justified while API intake and legacy fallback coexist	Retained shadows may eventually be deduplicated only after collector/OMC references, hashes and retention are resolved; no manual moving
One current ias-linkedin namespace	Justified: isolates explicit production paths from old flat wildcard/evaluation files	A separate folder per processing stage is not currently required; filenames/proofs express stage boundaries
Working outputs versus completed archive copies	Justified: permits reruns without replacing reviewed evidence	Keep copy/index distinction; choose expiry/storage policy separately
Reviews versus generated files	Justified: protects hashes and keeps human decisions distinct from automated proof	A future approval consumer would need its own schema and explicit trigger; none exists now
Synthetic acceptance versus real-privacy evaluation	Justified: synthetic checks and sensitive real-source evaluation have different privacy needs	Preserve isolation even if operator tooling is unified
Separate orchestration fixtures	Its isolation is justified; the standalone top-level name reflects the historical test implementation	Could join a managed test namespace after references/evidence are inventoried
archive-check, archive-check-final, archive-check-release	Historical iterations, not three documented enduring design roles	Candidates for a documented validation-evidence retention scheme; not cleanup targets today
Old flat outputs, Qwen quality variants and model traces	Historical leftovers from previous/evaluation workflows, not current handoffs	Classify provenance, sensitivity and consumers before any relocation; traces are not automatically public-safe
omc-krisp compatibility helper	Existing inactive-workflow reference explains its retention	Could be retired with a deliberate legacy rollback-policy change; it is not the active selector
Old meeting-named source folders	Retained imports outside current selector; exact original writer/rationale is unrecovered	No invented processing role or implied deletion permission

Separation or leftover	Assessment	Future consolidation consideration
Helper code mixed with current documents	Functional existing deployment choice: commands resolve helper paths under the mounted output area	A service-only code mount may be cleaner later, but would require deployment/path changes; none made here
model-eval filenames and inactive workflow display names	Compatibility/history, not accurate user-facing state labels	Renaming could improve clarity, but requires updating regexes, handoffs and rollback references together

No folder was created, moved, merged or removed for this documentation task. Only guide files and documentation links are deployed into the existing starting location.

Live state and reconciled records

The live instance has three published stages despite “inactive” in their display names. San is the sole enabled content clock: **Monday 07:00 America/New_York**, cron 0 7 * * 1. Enr's stored Monday 09:00 and Plan's Monday 10:00 schedules are disabled/disconnected; waiting parent calls invoke them. The next scheduled start is October 12, 2026 at 11:00 UTC. LinkedIn posting remains disabled.

Workflow ID	Published version observed October 7	Role
IASLinkedinSan01	0f5efe7d-b8ff-4547-8910-35ab7ecc1049	Selector, extraction/privacy approval and waiting enrichment call
IASLinkedinEnr01	70e431c0-d9b7-4a0d-a3b6-576b4d0ec7ac	Approved themes, Brave research, candidates and waiting planner call
IASLinkedinPlan1	acfd5d24-452d-42aa-af06-96cabe56b6b3	Editorial plan and verified archive/index completion

Earlier repository statements that all IAS copies are inactive describe migration preparation, not this deployed state. Git JSON templates still intentionally have inactive/disconnected clocks; they are not runtime-state exports. The original OMC-krisp selector reference belongs to the preserved inactive sanitizer, while live San uses /data/output/ias-linkedln/select-transcripts.cjs. The top-level activation snapshot in content-pipeline-state.json records the original planner publication; its later contentUsabilityRepair section records the current version above. The output date and UTC archive date are both correct. The earlier repair prose described the approval as already expired; its live timestamp instead sets the expiry to October 8 at 01:49:45.270 UTC. Its stored state can remain “approved” after that time while the helper rejects it as stale.

Repository references below support the current trace and distinguish historical records. Live mounts, workflow exports, helper behavior, current index/manifest, directory permissions and Samba share configuration take precedence when older prose conflicts.

- [Applied content access and archive record](#)
- [Sanitized focused acceptance](#)
- [Current state and subsequent repair metadata](#)
- [Transcript intake and identity repair](#)
- [Cutover history and applied publication](#)
- [Archive implementation](#)