

Content evidence completion — October 7, 2026

This follow-up replaces PR #18's prompt-encoding workaround with safe request construction, repairs observed page extraction and qualification defects, and reviews the five deferred topics against retained evidence. The n8n expression engine itself is unchanged. Production models, the Monday 07:00 America/New_York schedule and disabled LinkedIn publishing are unchanged.

Uninterrupted cycle result and review locations

The full sanitizer → privacy → research → page retrieval → assessment → planner → archive chain ran **18:07:49-18:13:20 UTC on October 7, 2026**, with CLI exit 0 and root execution status success. No checkpoint continuation, manual grouping or child-stage restart was used in this successful run. Eighteen URLs yielded **14 usable pages and four unavailable sources**; seven candidate groups produced **two evidence-bound opportunities and five deferrals**. This is a different input set from PR #18, not a controlled before/after quality score.

The accepted material provides **narrow factual starting points, not publishable posts**:

1. **NIST digital identity transition.** The [Revision 4 index](#) records the July 2025 final release (s6/p0002); the [Revision 3 index](#) records supersession as of August 1, 2025 (s13/p0001). Both claims passed the explicit attributed-document exception. Both pages are one NIST origin. A useful angle is checking whether an identity roadmap still relies on the superseded revision. Specific implementation requirements need their own evidence review; this is not new October news.
2. **Critical-infrastructure resilience policy context.** [CISA's overview](#) records the April 30, 2024 national security memorandum (s5/p0006). Only that attributed historical fact passed. Leadership implications, operational requirements and claimed benefits are not established by this one approved claim. This is the weaker editorial opportunity and needs a more concrete audience decision before drafting.

All five files were copied to the existing archive and verified against their hashes. All **12 prior archive files** remained unchanged. The exact retrieval packet is preserved at Evidence/packet-438a03d63be65f1de5eed620200a60264a7822f54f679a5dd5373a74a9284ef6.json. Original and re-extracted source snapshots remain available. Optional tests and model calls stopped after the completed cycle; subsequent work was limited to integrity checks, documentation, access checks and Git completion.

Open \\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedln\Evidence-Completion.html for this report, or the neighboring Evidence-Completion.pdf. Start editorial review at \\Iggy-Nas\Shared\ContentPipeline\output\ias-linkedln\CURRENT.md.

The immutable result is in:

\\Iggy-Nas\Shared\ContentPipeline\output\Archive\ias-linkedln\2026-10-07-bb7d4a34a9ddd7cb0dc60fa23fb4a0ee9b58c6000916e83efc9f226b70e22ddf

- [Editorial plan](../Archive/ias-linkedin/2026-10-07-bb7d4a34a9ddd7cb0dc60fa23fb4a0ee9b58c6000916e83efc9f226b70e22ddf/linkedin-editorial-plan-model-eval-2026-10-07.md): read the approved factual claims and sources before considering the proposed angle.
- [Research brief](../Archive/ias-linkedin/2026-10-07-bb7d4a34a9ddd7cb0dc60fa23fb4a0ee9b58c6000916e83efc9f226b70e22ddf/content-brief-model-eval-2026-10-07.md): accepted and deferred claims with passages and judgments.
- [Candidate/evidence record](../Archive/ias-linkedin/2026-10-07-bb7d4a34a9ddd7cb0dc60fa23fb4a0ee9b58c6000916e83efc9f226b70e22ddf/content-candidates-model-eval-2026-10-07.json): exact source bindings and gate reasons.
- [Plan review flags](../Archive/ias-linkedin/2026-10-07-bb7d4a34a9ddd7cb0dc60fa23fb4a0ee9b58c6000916e83efc9f226b70e22ddf/linkedin-editorial-plan-model-eval-2026-10-07.json): inspect approvedClaims, evidenceJudgments, fullProseReview and webinarReview.

Save edits and decisions in \\Iggy-Nas\\Shared\\ContentPipeline\\output\\ias-linkedin\\Reviews, identifying cycle bb7d4a34... and the topic. Leave generated and archived evidence unchanged. Review records are human records only; automation does not consume them.

Remaining deferrals in the new cycle

These are the new cycle's five deferred groups, distinct from the five historical groups reviewed below.

Group	Exact recorded block	Interpretation and next editorial action
NetScaler zero-day / SAML disruption (t1)	no_bound_material_claims, model_deferred	Poor topic/source selection remains. Its three retrieved sources concern Southern Company and Advantest incidents, not the NetScaler topic. Available text cannot establish that topic. Do not use the utility notice as NetScaler evidence. No extra retrieval was used to rescue the count.
Agentic AI governance / identity (t2)	no_bound_material_claims, model_deferred	Unavailable text plus weak topical evidence. Akamai and Dark Reading returned HTTP 403. The available Palo Alto page concerns vulnerability discovery rather than the proposed identity-controls claim, and its host is not in the approved primary registry. No bound claim was returned. Availability, scope and authority need resolution; the model's unrecorded rationale must not be invented.

Group	Exact recorded block	Interpretation and next editorial action
Critical-infrastructure isolation / recovery (t3)	independent_corroboration_unresolved on the first claim	Model contract failure with usable evidence. The model called a paraphrase a direct quote and requested the exact-quotation exception. Its text differs from s3/p0004, so that claim does not qualify for the exception and has only one origin. A second exact attributed CISA claim at s10/p0029 passed, but the candidate remains deferred because every proposed material claim must pass. A human can use the exact bounded guidance; this is not a finding that public evidence is absent.
Logging architecture / visibility (t4)	independent_corroboration_unresolved	Model contract failure with usable evidence. The model shortened/reworded CISA's release sentence while calling it a direct quote. The exact attributed exception therefore did not apply; only one source origin was bound. CISA's release , s18/p0003, supports a bounded document-release topic with the FCEB-agency and OMB qualifiers preserved. Do not broaden it into a mandate for all organizations. No new fetch is needed for human review.
AI governance / policy frameworks (t7)	no_bound_material_claims, model_deferred	Mixed retrieval and authority coverage limits. The NIST framework page produced unusable_page_text; the saved failure does not distinguish a challenge/empty page from an unhandled layout. The AI RMF Playbook page was retrieved, but airc.nist.gov is absent from the explicit primary/exception registry and the model returned no claims. This is not proof that the topic lacks public evidence. The host needs a deliberate authority review before an automated exception.

The result demonstrates that supported material is no longer rejected indiscriminately and that the parser path works end to end. It does **not** establish good topic discovery or ready-to-use writing. The accepted foundation facts are weakly connected to their incoming theme labels; topic relevance is still a model judgment rather than a deterministic guarantee.

The planner also retains generic fallback wording such as “verified guidance” and “independent coverage,” although these accepted claims use single-origin exceptions. Treat those phrases as editing defects, not evidence of independent verification. Both topics require webinar review and complete prose review. The current fullProseReview contains two coarse serialized-proposal blocks per topic; this is not a count of two factual errors or a sentence-level measure, and the final fallback prose is not separately scored. Robert and Leigh must review the entire final plan, including industry assignments and service claims. These existing drafting limitations were documented rather than expanded into another production repair or evaluation run.

What was fixed

The original inline expression failed when the prompt contained nested closing braces. PR #18 encoded the prompt to get past that parser; its continuation also reconstructed topic groups by source-ID order because the exact packet was missing. That continuation did not demonstrate an uninterrupted pipeline.

The Code node now builds a normal request object, with prompts and evidence held as data. The HTTP node evaluates only `={{ $json.requestBody }}`. The focused offline check reproduced the original expression failure in the installed n8n evaluator and proved that the new construction preserves braces and quotes. This fixes the pipeline's underlying request-construction defect without a vendor-engine patch or service restart.

The extractor now selects Microsoft article prose instead of related cards and recognizes CISA's paragraph-free text containers. A generic canonical homepage cannot replace a document URL. Old successful cache entries are re-extracted from saved HTML into new immutable snapshots; original snapshots remain. The shortlist keeps actual topic groups and allocates its existing 18-URL budget across groups rather than exhausting it on the earliest groups.

Qualification normalizes whitespace, Unicode and typographic quotes without substituting words. An exact official quotation explicitly requesting the existing exception can acquire its required publisher attribution; paraphrases cannot. The reviewed authority registry now includes the Canadian Cyber Centre, National Laboratory of the Rockies and Southern Company for their own document statements. News reporting does not become primary evidence, and unknown or shared origins do not become independent corroboration.

Review of PR #18's five deferrals

These findings concern the saved five groups, not a replayed model evaluation. Public links identify sources; full retained text and passage bindings stay in the shared Evidence cache and private checkpoints. Proposed reframings below are editorial opportunities requiring human review, not automatically approved drafts.

Saved topic	Actual blocker and diagnosis	Supported scope / remaining deferral
Georgia Power / Alabama Power, 400,000 accounts	Source independence and unsupported quantitative claims. News was incorrectly labeled primary; unrelated stories had been grouped with it during recovery. Quote typography was a secondary binding defect.	The company notice supports a bounded, explicitly attributed account of the information involved and the company's response. It does not establish the 400,000 total or 300,000/100,000 breakdown. Repetition of the notice is one origin, not corroboration. The broad numerical headline remains deferred.
Data Centers Gap Analysis	Poor topic selection and source-role/independence judgments. NLR research was labeled independent reporting; the recovered group mixed US research, Finnish forecasts and European grid reporting.	NLR's retained research abstract and modeling research support attributed discussion of AI workload variability, grid modeling and siting trade-offs. Two NLR publications are not independent origins. Broad cross-geography forecasts remain unsupported; a bounded research-document topic is viable. No extra retrieval was needed.

Saved topic	Actual blocker and diagnosis	Supported scope / remaining deferral
Microsoft Digital Defense Report 2026	Incorrectly applied qualification format; a separate extraction defect. The saved response explicitly requested the exception with exact source text but omitted the literal publisher prefix. That alone should not lose an otherwise eligible quote.	The report supports attributed statements about Microsoft's observations and survey figures, preserving population, dates and methodology. An offline replay of the saved response now accepts the exact attributed statements. This is not independent confirmation of Microsoft's claims. The linked RMM article's apparent unavailability was an extractor defect, now repaired from saved HTML.
Canadian CIREN isolation guidance	Incorrect exception eligibility and claim scope. The government authority was eligible as primary but absent from the narrower exception registry; paraphrases and typography also failed exact binding.	Canadian guidance supports attributed contingency planning for essential operators. Preserve the last-resort condition, locally defined industry triggers and “up to three months” scope. It is not a universal isolation mandate. The bounded quotation passes the offline check; an invented universal 9,000-day claim fails. Joint guidance remains one origin even with multiple agency logos.
NIST Digital Identity Guidelines	Poor selection and a canonical-URL defect. The old index's generic root canonical erased its document identity; “documents are available” was a weak editorial angle.	Revision 4 , linked by the saved legacy index, supports an attributed current-guidance transition topic. Its final release is July 2025, not a new October 2026 event. Do not claim implementation compliance or outcomes from the index alone. The unrelated CISA CDM PDF was not pursued to fill a source count.

The four targeted fetches were the linked Southern Company notice, linked NIST Revision 4 index, the previously unusable [Microsoft RMM article](#), and [CISA's isolation resource](#). All four yielded usable page text with extractor 1.1. The latter two were re-extracted from the same fetched HTML after correcting extraction. No broad extra search or PDF collection was used for this review.

Acceptance and operating limits

The focused completion check reused saved evidence without network or model calls. It reproduced the parser defect, exercised the corrected request, accepted valid attributed Microsoft and Canadian material, retained the uncorroborated incident-count deferral, rejected an exaggerated universal claim, recovered the two page bodies and preserved document canonical identity. PR #18's 15-case gate is reused rather than repeated. These checks establish mechanics and policy behavior with saved or stipulated judgments, not semantic-model accuracy.

The first attempted full run stopped at research privacy review before outbound research. Its private checkpoint is retained. The old error omitted rejection reasons, so the precise rejected concept cannot be reconstructed; the rejection is not evidence of a parser failure. Query instructions now explicitly match the existing source-supported privacy contract, and future failures identify rejected indices and risk categories without exposing query text. The privacy guard still fails closed and was not bypassed.

Human review remains necessary for meaning, source-origin judgments, dates, qualifiers, editorial relevance, webinar overlap and unmatched prose. Model confidence and same-model review are not independent validation. Human editing effort and alternative-writer quality remain unmeasured. Source text, transcripts, credentials and model evaluation packets are excluded from Git. There is no automatic publishing or new production drafting stage.